

# WHAT HAPPENS WHEN YOU TYPE IN A URL & HIT ENTER

*Way more detail than you could possibly ever want, guaranteed!*

Presented by Nick Evans for ECT, June 2024

# WHY THE QUESTION?

- Touches EVERYTHING in Tech
  - o No matter your role - some (probably most) parts of this impact your job day-to-day
  - o Applicable if you're service desk, web dev, analyst, PM, telecom tech, or even a cabling contractor
- Great for interviews
  - o Open-ended
  - o Can go in a lot of directions
  - o Which part(s) people focus on can be interesting 🤔
- Kind of a fun deck: lots of gotchas and little details you don't generally think about!

WHAT HAPPENS WHEN YOU  
TYPE IN A URL & HIT ENTER

# WELL, WHAT KIND OF KEYBOARD ARE YOU USING?

- You're *probably* using a USB keyboard, so let's assume that for the sake of brevity.
  - **BUT:** works a bit differently with an olde-style PS2 keyboard!
  - Directly sends interrupt requests to the processor on IRQ1



- But before we get into how USB keyboards work, consider the physical device!
- Keycap -> switch -> matrix circuit doing “the magics”
  - Switches are a whole thing on their own: membrane vs. mechanical
  - Inside mechanical switches, there's linear vs. tactile
  - People optimize their switch type for task -- e.g. linear offers the shortest travel time (at the cost of no tactile feedback)

# USB KEYBOARD: CRASH COURSE

- USB spec defines standard for the most common HIDs: keyboards, mice, and game pads
- Plugging your keyboard in causes an intense negotiation to register it w/ the USB host & potentially ask it for more electricity
  - To power your RGB, of course
- Pressing a key has distinct states: key down, key up
- Keys have integer key codes
- When you press something, the keyboard's own hardware registers it & stores that in its own buffer
  - Key down 13, key up 13 (the enter key!)
- USB host controller polls keyboard every **X** ms
  - Polling period set up during negotiation!

# KEY PRESS DATA GETTING TO SOFTWARE

- USB host controller gets the data out of the keyboard's buffer
- Hands raw data off to the USB HID Keyboard driver (which is generic)
- USB HID driver passes keypress data to the OS' hardware abstraction layer
- Of course, for non-USB keyboards, some different stuff happens!
- But it all eventually ends up hitting the OS' hardware abstraction layer
  - Which can then hand the keypress events off to its software
  - So devs don't have to care what kind of keyboard it is!

# DE-RAIL: KEYBOARD HARDWARE LIMITATIONS

- Physical keyboards occasionally are bad at detecting keypresses when you're mashing lots of buttons!
- This usually happens with cheaper keyboards: the circuit isn't really designed with more than 1 or 2 keys being pressed at once in mind
- So the keyboard matrix circuit may mis-detect a combination of 3 keys as 4 keys
  - And then another anti-ghosting measure will detect \*that\* and drop the "ghost" 4th key press
- Further fun: not every circuit behaves the same, so you end up with different combinations causing problems across different models!
- Higher-end keyboards feature 3-key (or N-key) rollover
  - Basically, they spent more money on their keyboard matrix circuit and it scans for key presses on each individual button
- Microsoft has a demo of the ghosting & N-key rollover detecting stuff



# DE-RAIL: OTHER TYPES OF KEYBOARDS

- PS2 & USB keyboards aren't the only types available though.
- Perhaps the most common in 2021: virtual keyboards!
  - Cuz, you know, iOS/Android
- Apple touch screens feature capacitive layer that detect electrical changes where you put your fingers
  - Keydown / keyup events are detected in software by the screen reporting touch event  $X$ ,  $Y$  coordinates
- Much more reliant on software (obviously!)



# COOL, KEY PRESSED!

We're seven slides in  
and all we've talked about is keyboards 😊

# OS ROUTES THE KEYBOARD EVENTS

- Oh, did you think we were done talking about keys?
  - We've got keyboard input into the OS
  - But it still needs to find its way into a userland application like Chrome
- OS is aware of which application has focus
  - Differs a bit between Windows/UNIX-likes
  - Win32K (the kernel) is responsible for giving focused userland app the keyboard events
  - On Linux, the kernel just barfs keypresses and a userland GUI server like X.org is the intermediary between the kernel & application
    - MacOS is similar to Linux

# APPLICATION RECEIVES A KEY PRESS

- Applications receive keypresses as events and generally have event listeners bound
- When you focus a text input field like the URL bar, your GUI API generally helps you out and just fills it in as people type
- But apps can also listen for specific events w/out an input field
  - e.g. ctrl+r = reload the page
- Let's assume "google.com" has gone into Chrome's URL bar
- but we **haven't** hit enter yet!

# THE CHROME URL BAR

- The term “URL bar” isn’t a very good description, ‘cuz it does a couple things when you type into it:
  1. Check your history
  2. Check preferred search engine for potential hits & related search terms
  3. Look stuff up in Wikipedia & other fact-giving services
- That’s all before even hitting *<enter>*
- Checking your history consults a sqlite database Chrome has in your profile
  - It factors in a couple things, including how recently you’ve gone to that site
- But the suggested search terms & looking up celebrity photos from Wikipedia/etc goes over the network to Google’s APIs

# NETWORKING, PART 1

This is going to be dense, since I have to cram 60 years of incredibly boring history into a slide deck.

*Might wanna grab another slice of pizza at this point ...*

# OVER THE NETWORK AND THROUGH THE WOODS TO G'S HOUSE WE

## GO

- Talking to the network is an involved process
  - But we're getting ahead of ourselves!
- For brevity, we'll assume the network is already set up and is able to talk to the internet
  - So we won't go into stuff about IP assignment via DHCP or otherwise
- Your computer has a gateway to the rest of the internet, and knows to route traffic over to that gateway
  - Chrome is going to want to hit some `<google.com>` URL to get the suggestion data
    - But `<google.com>` isn't an IP that your computer can route to, so you have to do a DNS lookup.

# WELCOME TO DNS

- We invented DNS to turn nice domain names into IP addresses after the internet got too big to remember everything
- It's a distributed fault-tolerant system
  - Which often is at fault for outages, ironically
  - But this isn't an ops slide deck so we won't make too many jokes at DNS' expense
- DNS servers get requests from clients: "resolve google.com please"
- Your closest DNS server (probably run by your ISP) isn't in charge of <google.com>, so it has to figure out who is in charge of that.
- So DNS servers ask a chain of increasingly-important servers what IP google.com resolves to

# RESOLVING A DOMAIN NAME

- Your system's configured resolver (generally called a *recursive resolver*) will ask a root nameserver about `<.com>`'s authoritative nameserver.
  - The root nameservers are the 13 "servers" that make the whole internet's work
- Root nameserver tells you to ask some IP about `<.com>` domains
- Recursive resolver then asks the `<.com>` server what server is in charge of `<google.com>`
- The recursive resolver can then ask it for the IP that `<google.com>` has
- If you had more subdomains, this could continue for more levels
  - E.g. `<api.chrome.google.com>`



# DNS RECORDS

- Of course, DNS is more complicated than “what’s google.com’s IP address?”
- DNS records you’d be looking for here are A & AAAA
  - But if you had `<chrome.google.com>`, you could be looking for A, AAAA, CNAME, or ALIAS.
- The A & AAAA will give you IPv4 & IPv6 addresses, respectively.
- CNAMEs and ALIASes are pointers to other records:
  - `<chrome.google.com>` could have an A record for 8.8.8.8, or a CNAME pointing you to `<browser.google.com>`
    - So CNAME = more recursion!

# DOWN THE RABBIT HOLE: NETWORKING

- We've already used the IP protocol to make DNS requests, but I haven't explained what's happening there!
- The IP protocol is a means of addressing packets of data to other networks, using IP addresses.
- IP protocol works irrespective of the physical transport (ethernet, fiber, etc) & irrespective of what the packets of data contain (e.g. Netflix or timesheets)
- IP is essentially "the internet" -
  - this is a logical layer that everything talks to each other with
- IPv4 and IPv6 serve the same purpose, but use different addressing schemes
  - IPv4 is limited to 4.2b IPs, since they used a 32-bit integer for the IP address field in the spec
  - But that's NOT ENOUGH
  - IPv6 uses a 128-bit integer, which is enough IPs for every single molecule in the universe

# BUT HOW DOES MY COMPUTER ACTUALLY MOVE IP PACKETS?

- That depends!
- The other “layers” above & below IP are meant to be modular
  - So you can send IP packets over dial-up or fiber optic cable
- Generally, your computer will have an Ethernet card with a copper wire hooked up to a router.
  - Wifi will pretend to work the same, except with radios instead of wires. And some additional error-correction done below the IP protocol.
- When you want to send a packet to Google via your router, your computer will hand it off to your router’s IP
  - So your Ethernet card has to make an ARP request on the network, asking if anyone knows the MAC address for your router’s IP
  - E.g. “hey what’s the ethernet card I’m connected to that has 192.168.0.1 as its IP?”
  - And the router will see this shouted out onto the network and say “It’s me! Send it to 00:0a:95:9d:68:16!”

# BUT HOW DOES ETHERNET ACTUALLY MOVE PACKETS?

- At this level, they aren't considered "packets". That's a concept for the IP protocol.
  - We call 'em frames down here, and they can be sliced up as the hardware requires
- But ethernet uses electrical signals over copper!
  - Fiber will use light through glass
  - WiFi will use extremely shrill noises over air
- The hardware will address other stuff physically connected to it (e.g. MAC addressing for Ethernet)
- And translate some kind of signal into data before passing it back up the stack (or vise-versa)

# OK, BUT HOW DO I KNOW MY PACKETS GOT THERE?

- Above IP, we can have the TCP protocol
  - TCP is all about transmission control
  - You shake hands and acknowledge receiving data
  - So Chrome knows its packets sent to Google's servers are being received!
- But we got here from DNS, so ignore everything I just said.
- DNS requests don't (typically) use TCP -- it uses UDP instead!
- TCP adds a whole bunch of overhead for what is otherwise a very small bit of data: "what is the IP for <x>" / "the IP is <x>"
- That response is acknowledgement that the request was received, so that part of TCP is useless too

# OK, WELL, HOW DOES MY COMPUTER DEAL WITH TCP OR UDP?

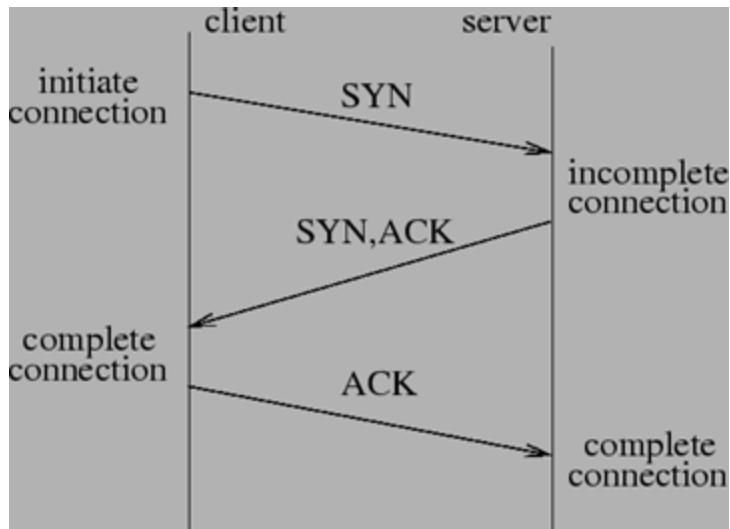
- Your kernel has a network stack that knows about UDP & TCP, IP, ARP, and whatever you need for your hardware.
- Developers generally don't have to care about this, beyond indicating if they want TCP or UDP when they open a socket
- “Wait,” you say, “*what’s a socket?*”
  - A socket is an open connection to an IP address on a port using a transport layer protocol
  - It can send data back & forth between the two computers
- “Port?”
  - It’s an apartment number to the IP’s address.
  - A server will let its various services (website, DNS, email) listen to a port
  - So when you connect to port 53, you’re *probably* talking to a DNS server

# MAKING THE API REQUEST TO GOOGLE

- Getting back to the topic at hand...
- Once your recursive resolver gets the IP address, it will cache it for some time
- And now Chrome can ask Google if there are any suggested search terms, good search results, or pictures of Nicholas Cage to display as potential options!
- This API request happens over HTTPS
  - Chrome builds a URL, e.g.  
`https://suggestions.google.com?k=<what you typed>`
- Chrome opens a TCP socket to the IP on port 443

# REVISITING TCP

- We brushed off TCP earlier since we were worried about DNS requests over UDP
- But a TCP request is a much more involved process than “send a packet, hope it arrives”
- Everything requires confirmation & confirmation of getting the confirmation
  - Opening the socket
  - Each packet being sent
  - Closing the socket





# TCP IS NOISY

- With all the back & forth acknowledgments, TCP requires a lot more bandwidth
  - And if something isn't acknowledged properly, it will try to re-send the packet
- TCP packets also contain information about their order
  - So the receiver can get them out-of-order and hold them in a buffer until it's got the missing packets
- The tradeoff between TCP and UDP is reliability: you can guarantee stuff is getting there in the right order with TCP
- Which makes it great for HTTP, since you're sending so much more data than “what is <x>” / “it is <x>”

# SETTING UP THE HTTPS SECURE TUNNEL

- HTTPS is the good-ol' HTTP protocol sent through an encrypted TLS tunnel
- Setting up the TLS tunnel happens first
- TLS guarantees two things:
  - The connection is free from eavesdroppers (**encryption**)
  - The server belongs to whomever the certificate says it does (**trust**)
- The server will proffer a certificate with information about who issued it, what website they issued it for, when they issued it, and how long it's valid for
- Browser checks all of this info
  - Current time is between issued at & expiration dates
  - Certificate is for the domain name Chrome is talking to
  - Certificate was issued by a trusted Certificate Authority

# CERTIFICATE AUTHORITIES

- Chrome has a list of Certificate Authority certs that it trusts
  - This list is maintained by a consortium: Google, Mozilla, Apple, Microsoft, etc all sit on the board
- These CA certs sign a website's cert
  - The CAs are *supposed to* verify that the person paying them \$15 for a cert on northwestern.edu is actually part of Northwestern University
  - This is how trust is established!
- The website's cert may not be signed directly by a trusted CA key
- Many CAs use one (or more) "intermediate" certificates that they can rotate out more easily than the trusted CA cert that *every single browser bundles & ships to millions of computers*
- Webserver may need to help the client out by proffering the whole chain of intermediate certificates

# DE-RAIL: NAME-BASED VIRTUAL HOSTING & TLS CERTIFICATES

- Name-based virtual hosting is the practice of having one webserver on one IP serve several websites
- The server knows which site to give you based on the HTTP *Host: google.com* header
- But the TLS tunnel is established *before* the HTTP request is made
- So the TLS folks came up with the Server Name Identification (SNI) extension to TLS
- That *Host: google.com* header becomes part of the TLS handshake, so the webserver can look up the right certificate to proffer
- This is an OPTIONAL extension to TLS
  - It's widely supported now, but real old WinXP machines may have issues

# REVOKING CERTIFICATES

- A website's certificate consists of two parts:
  - Public key: contains all the data we've been talking about)
  - Private key: used to encrypt data
- The private key is intended to remain private. If it is leaked, somebody could decrypt the traffic and use it for evil!
- Accidents (and Heartbleed) happen, so there is a way for a CA to revoke a certificate
  - Actually several ways...
- When a browser is proffered a certificate, it will check two things to see if it's been revoked:
  - Certificate Revocation List
  - OCSP Status

# CERTIFICATE REVOCATION METHODS

- Certificate Revocation Lists are a big list of certificates that have been revoked
- Browser downloads this list every so often, and then it can check it (ON YOUR COMPUTER) whenever it needs to
- Everyone downloads the full list
- OCSP is more like an API that you ask about a specific certificate
- Advantage over CRL is that the OCSP status is *always* correct, whereas your CRL may not have been updated in the last few hours
- Disadvantage is that somebody may have a record of what specific site you're visiting

# CERTIFICATE REVOCATION METHODS: PART 2

- Not every client uses CRLs or OCSP
- But browsers all use *at least* one method, if not both.
- Certificate revocation may not be in sync between the CRL and OCSP immediately

- The long and short is:

Troubleshooting a “bad cert”  
when it’s been revoked is  
***absolutely maddening***

# CHECKING IN ON DNS

- That privacy concern from OCSP a few slides back is a good segue!
- DNS is not using TLS. It was invented long before we had enough processing power for encryption.
- This is bad for two reasons:
  1. How can I trust DNS responses aren't being tampered with by my ISP?
  1. How do I know I can trust the DNS server?



# DNS SECURITY FEATURES

- DNSSEC is an optional extension for DNS that adds public/private key encryption to records.
- An additional record that your recursive resolver can ask for (*DNSKEY*) is available
  - This is a public key you can use to validate a hash that comes back with your A/AAAA/etc records
- Every server in the chain that your recursive resolver follows must be signed
  - This is similar to the CA cert signing a website's cert
- But there's no inherent encryption
  - Chrome & Firefox are experimenting with "DNS over HTTP" so they can take advantage of TLS

# THE HTTP PROTOCOL

We're 32 slides in and finally getting to an *HTTP GET*.

# ANATOMY OF AN HTTP REQUEST

HTTP GET /suggest?v=google.com HTTP/1.1

Host: google.com

User-Agent: ChromeLongStringBlabla 77 (Mozilla/IE/LOLCATS)

Accept: \*/\*

{no body}

# ANATOMY OF AN HTTP RESPONSE

HTTP/1.1 200

Content-Type: application/json

Expires: Wed, 25 Jan 2021 15:15:22 GMT

Content-Length: 220

```
{“suggestions”: [{url: “google.com”, ...}]}
```

# COMPRESSION

- HTTP supports negotiation via headers
- Client can specify **Accept-Encoding: gzip** (or a couple other values)
- Server will compress the response & include a **Content-Encoding: gzip** header to indicate that it supported the compression requested
- Compressing the response saves on bandwidth, at the cost of a small amount of CPU client-side

# HELLO, TCP

Remember we said TCP is noisy and unnecessary for DNS?

Well, we use it for HTTP because the packets that make up a request/response need to be in order!

# CHROME GETS SUGGESTIONS BACK & RENDERS THEM

- API call complete, Chrome now has the data it needs to build the suggestions!
- Parse the JSON from the API call, merge it together w/ history data
- We can finally hit `<enter>`!
  - For brevity, let's assume we didn't pick one of the suggestions

# CHROME EVALUATES THE “URL”

- Again: “URL bar” is a misnomer
- Chrome has to check to see if you’ve entered a valid URL
  - Otherwise it will send you to Google’s search results
- “URL” evaluation is fairly involved
- Does it have a protocol?
  - If not, assume https
  - If it does, is that protocol supported? E.g. *ftp://* was dropped in Chrome 82
- So our `<google.com>` should turn into `<https://google.com>`



# SAFE BROWSING CHECKS

- Chrome (and most other browsers) have a “safe browsing” list
- Offers the Safe Browsing Update List API
  - Browser sends a hash on a partial URL
  - Response indicates if it's evil
- Privacy implications:
  - You tell Google every site you visit
- Google maintains this list, but there are alternatives
- When the URL is determined to be **Evil**:
  - Chrome shows a warning screen instead of proceeding w/ HTTP request
- But, `<google.com>` won't trigger a warning.

# NOW WE REPEAT STUFF

The DNS lookup, TCP/IP connection, TLS negotiation, revocation checks, and HTTP request/response all take place again.

This time, for your URL instead of the suggestion API.

# DE-RAIL: INTERNATIONAL DOMAINS

- DNS requires names to be ASCII
  - aka: the English alphabet
- But not everyone speaks a language that works w/ our alphabet
- Option 1: update every DNS server to support unicode hostnames
  - Internet says: no
- Option 2: punycode!
- This is a way for unicode characters to be represented as ASCII in DNS
- *<Bücher.example>* - must be punycoded!
  - Prefix punycode with *xn--*
  - Thus: *<xn--bcher-kva.example>*

# DE-RAIL: THE CACHE

- I've assumed that Chrome *doesn't* have the google homepage cached.
- Before the browser makes a connection, it will consult its local cache to see if it has that page already & can skip all the network stuff.
- Depending on headers on original response, browser may need to re-validate cached page before it can use it
- May send an HTTP OPTIONS request to check the Last-Modified or ETag headers
- These headers can be used by the server to indicate if the page has changed
- If it hasn't changed: use cached copy!

# BETWEEN REQUEST & RESPONSE

We looked at an HTTP request and an HTTP response.

But what happens between them?

# BETWEEN THE REQUEST & RESPONSE

- Short answer:
  - A webserver (eventually) gets the request
  - Builds a response: static files, or code runs
  - And then it builds the response
- Actual answer is complicated and varied
  - There are many, many, many ways to serve web traffic.
- In the case of the Google homepage...
  - The domain or IP are *probably* being served by your closest datacenter
  - *Probably* hits a load balancer
  - And the homepage of Google.com *probably* comes from a cache on their side
    - Partially: your login info is still there!

# BETTER EXAMPLE (FOR OUR PURPOSES)

- A better example may be looking at something more familiar!
  - How about a request served from NU's AWS account by Lambda?
- Traffic gets to AWS over the internet
  - Hits the API Gateway in a region
  - Or hits CloudFront and is carried via optimized network paths back to the region, for Edge-Optimized API Gateways
- API Gateway unpacks the HTTP request & converts it into an event
- Is configured to invoke a Lambda
- Lambda spins up, looks at the event, runs code, builds a string response, and feeds it back through the chain to the client

# A MORE TRADITIONAL EXAMPLE

- Or in the case of NU's on-prem stuff, which is largely Apache ...
- Request comes in to the server (maybe through a LB)
- Apache looks at the Host header (or SNI info), finds the right virtual host config
- Might serve up a static asset
  - May hand off to PHP-FPM or mod\_php to build a response
- Might even go through the ForgeRock SSO module for Apache before processing the request
  - Which is looking at cookies & making more requests to back-end SSO servers



# RENDERING A WEBPAGE

We're getting there!

# RENDERING THE DOCUMENT

- Chrome looks at the *Content-Type* header in the response to figure out what the document even is
- *text/html* is what we're mainly concerned with
  - But it can deal with other things, like *application/pdf* or *image/png*
- HTML gets parsed, turned into a DOM, and rendered
- HTML can contain references to other resources
  - Javascript files
  - Fonts
  - Images
  - Stylesheets
  - Videos
  - Etc etc etc

# FETCH OTHER RESOURCES

- Chrome will try to fetch linked resources as it encounters them
- Before it grabs something, it may need to check CORS to figure out if it *should* load the resource
- Caching rules apply here (and are usually more important here)
- JS will be fetched & run when it is encountered in the DOM
  - This means your JS can run before the HTML is fully processed & rendered!
- Same is true for CSS, but those rules will apply to any element matching their selector
  - Elements loaded or soon-to-be loaded!

# DE-RAIL: CROSS ORIGIN RESOURCE SHARING

- CORS is a security mechanism implemented **by browsers** to stop malicious folks from making HTTP requests “as your browser”
- If `<evil-nick.com>` could make a request through your browser, with your login info, to `<amazon.com/order-a-car>` ...
- So browsers restrict requests that cross origins (domain A -> B)
  - This is mainly applicable to XMLHttpRequests (AJAX)
  - But some resources have this as well, e.g. fonts & JS
- So when we're loading resources, the browser will examine the response's CORS headers to see if your origin (domain A) has been permitted to load stuff from this domain (domain B)

# SO HOW DOES HTML RENDERING & JS EXECUTION WORK?

- **Hah, nope, sorry**, that question is too ambitious even for this obnoxious over-the-top joke slide deck.
- Rendering a webpage in 2024 correctly and securely is the hardest problem that mankind has invented.
  - The moon? *Please, we did that with a TI-84.*
- Problem is so hard, Microsoft gave up on it & started using Chromium for Edge

# SO THAT'S IT

You now know what happens.

There's a lot more that could be discussed which wasn't relevant to the question, like cookies, POST, streaming, and more.

I HOPE THIS TALK HAS SHOWN YOU THE INCREDIBLE COMPLEXITY  
OF THE WORLD WE HAVE WROUGHT

IT IS YOUR JOB TO KEEP THE HOUSE OF CARDS STANDING FOR THE  
NEXT 40+ YEARS

*please enjoy your career in technology!*



SCAN THE QR CODE FOR A COPY OF THIS DECK